

Exploratory Data analysis - EDA

```
In [ ]: # importing the required libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings('ignore')
```

```
In [ ]: #changing the location to read the file
import os
os.getcwd()
```

```
Out[ ]: 'C:\\Users\\Jeeva46\\Desktop\\GIT'
```

```
In [ ]: os.chdir("C:\\Users\\Jeeva46\\Desktop\\GIT")
```

```
In [ ]: # Loading the file to df
df = pd.read_csv('laptop_price.csv', encoding='latin-1')
df.head(5)
```

```
Out[ ]:
```

	laptop_ID	Company	Product	TypeName	Inches	ScreenResolution	Cpu	Ram	Memory	
0	1	Apple	MacBook Pro	Ultrabook	13.3	IPS Panel Retina Display 2560x1600	Intel Core i5 2.3GHz	8GB	128GB SSD	Int Gra
1	2	Apple	Macbook Air	Ultrabook	13.3	1440x900	Intel Core i5 1.8GHz	8GB	128GB Flash Storage	Int Gra
2	3	HP	250 G6	Notebook	15.6	Full HD 1920x1080	Intel Core i5 7200U 2.5GHz	8GB	256GB SSD	Int Gra
3	4	Apple	MacBook Pro	Ultrabook	15.4	IPS Panel Retina Display 2880x1800	Intel Core i7 2.7GHz	16GB	512GB SSD	Ra Pr
4	5	Apple	MacBook Pro	Ultrabook	13.3	IPS Panel Retina Display 2560x1600	Intel Core i5 3.1GHz	8GB	256GB SSD	Int Gra

```
In [ ]: df.shape
```

```
Out[ ]: (1303, 13)
```

- knowing the dimensions of data
- 1303 samples and 13 features

```
In [ ]: df.duplicated().sum()
```

```
Out[ ]: 0
```

- No duplicates where found

```
In [ ]: df.isnull().sum()
```

```
Out[ ]: laptop_ID      0
Company      0
Product      0
TypeName     0
Inches       0
ScreenResolution 0
Cpu          0
Ram          0
Memory       0
Gpu          0
OpSys        0
Weight       0
Price_euros  0
dtype: int64
```

- No null values are present in dataset

```
In [ ]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1303 entries, 0 to 1302
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype
---  -
0   laptop_ID             1303 non-null  int64
1   Company               1303 non-null  object
2   Product               1303 non-null  object
3   TypeName              1303 non-null  object
4   Inches                1303 non-null  float64
5   ScreenResolution      1303 non-null  object
6   Cpu                   1303 non-null  object
7   Ram                   1303 non-null  object
8   Memory                1303 non-null  object
9   Gpu                   1303 non-null  object
10  OpSys                 1303 non-null  object
11  Weight                1303 non-null  object
12  Price_euros           1303 non-null  float64
dtypes: float64(2), int64(1), object(10)
memory usage: 132.5+ KB
```

- Getting info on the datatypes of the features

```
In [ ]: df.describe(include='all').T
```

```
Out[ ]:
```

	count	unique	top	freq	mean	std	min	25%	50%	
laptop_ID	1303.0	NaN	NaN	NaN	660.155794	381.172104	1.0	331.5	659.0	9
Company	1303	19	Dell	297	NaN	NaN	NaN	NaN	NaN	I

	count	unique	top	freq	mean	std	min	25%	50%	:
Product	1303	618	XPS 13	30	NaN	NaN	NaN	NaN	NaN	I
TypeName	1303	6	Notebook	727	NaN	NaN	NaN	NaN	NaN	I
Inches	1303.0	NaN	NaN	NaN	15.017191	1.426304	10.1	14.0	15.6	
ScreenResolution	1303	40	Full HD 1920x1080	507	NaN	NaN	NaN	NaN	NaN	I
Cpu	1303	118	Intel Core i5 7200U 2.5GHz	190	NaN	NaN	NaN	NaN	NaN	I
Ram	1303	9	8GB	619	NaN	NaN	NaN	NaN	NaN	I
Memory	1303	39	256GB SSD	412	NaN	NaN	NaN	NaN	NaN	I
Gpu	1303	110	Intel HD Graphics 620	281	NaN	NaN	NaN	NaN	NaN	I
OpSys	1303	9	Windows 10	1072	NaN	NaN	NaN	NaN	NaN	I
Weight	1303	179	2.2kg	121	NaN	NaN	NaN	NaN	NaN	I
Price_euros	1303.0	NaN	NaN	NaN	1123.686992	699.009043	174.0	599.0	977.0	148

In []: `df.nunique()`

Out[]: `laptop_ID` 1303
`Company` 19
`Product` 618
`TypeName` 6
`Inches` 18
`ScreenResolution` 40
`Cpu` 118
`Ram` 9
`Memory` 39
`Gpu` 110
`OpSys` 9
`Weight` 179
`Price_euros` 791
dtype: int64

- Reading the unique values
- This helps to know the countinuous variable and discrete variables
- Later we can use to enoding

In []: `df.drop(columns=['laptop_ID', 'Product'], inplace=True)`

- Dropping the unwanted features present in dataset
- These two feature are reduntant to analyse the price

Feature Engineering

Data visualization & feature engineering

Company column

In []:

```
# Create a figure with a specified size
plt.figure(figsize=(12, 10))


print('The no of unique companies: ',df['Company'].nunique())

# First subplot: Countplot
plt.subplot(2, 1, 1) # 1st subplot
sns.countplot(x='Company', data=df)
plt.xticks(rotation=90) # Rotate x Labels for better visibility
plt.title('Count of Laptops by Company')
plt.grid(True, axis='y', color='gray', linestyle='--', linewidth=0.5) # grids would

# Second subplot: Barplot
plt.subplot(2, 1, 2) # 2nd subplot
sns.barplot(x='Company', y='Price_euros', data=df,ci=None)
plt.xticks(rotation=90) # Rotate x Labels for better visibility
plt.title('Price of Laptops by Company')
plt.grid(True, axis='y', color='gray', linestyle='--', linewidth=0.5)

plt.tight_layout() # Adjust the layout to prevent overlap
plt.show()
```

The no of unique companies: 19



TypeName column


In []:

```
plt.figure(figsize=(12, 10))

# First subplot: Countplot
plt.subplot(2, 1, 1) # 1st subplot
sns.countplot(x='TypeName', data=df)
plt.xticks(rotation=90)
plt.title('Count of Laptops by Type')
plt.grid(True, axis='y', color='gray', linestyle='--', linewidth=0.5)

# Second subplot: Barplot
plt.subplot(2, 1, 2) # 2nd subplot
sns.barplot(x='TypeName', y='Price_euros', data=df, ci=None)
plt.xticks(rotation=90) # Rotate x Labels for better visibility
plt.title('Price of Laptops by Type')
plt.grid(True, axis='y', color='gray', linestyle='--', linewidth=0.5)

plt.tight_layout()
plt.show()
```



Inches column

In []:


```
plt.figure(figsize=(12, 10))

# First subplot: Countplot
plt.subplot(2, 1, 1)
sns.countplot(x='Inches', data=df)
```

```
plt.xticks(rotation=90)
plt.title('Distribution of Laptops by Size')
plt.grid(True, axis='y', color='gray', linestyle='--', linewidth=0.5)

# Second subplot: Scatter plot
plt.subplot(2, 1, 2)
sns.swarmplot(x='Inches', y='Price_euros', data=df)
plt.xticks(rotation=90)
plt.title('Price of Laptops by Screen Size (Inches)')
plt.grid(True, axis='y', color='gray', linestyle='--', linewidth=0.5)

plt.tight_layout()
plt.show()
```



Operating System OS

```
In [ ]: print(df['OpSys'].value_counts())
```

```
Windows 10      1072
No OS           66
Linux           62
Windows 7       45
Chrome OS       27
macOS           13
Mac OS X        8
Windows 10 S    8
Android         2
Name: OpSys, dtype: int64
```

- We can group the different categories of values
- Windows OS is distributed as more than one
- same like mac
- we can segregate to smaller bins

In []:

```

opsys = {
    'No OS': 'OS', # For null OS imputing OS
    'Android' : 'Android',
    'Chrome OS': 'Chrome',
    'Linux' : 'Linux',
    'Windows 10': 'Windows',
    'Windows 7': 'Windows',
    'Windows 10 S': 'Windows',
    'Windows S': 'Windows',
    'macOS': 'Mac',
    'Mac OS X': 'Mac'
}

for old_word,new_word in opsys.items():
    df['OpSys'] = df['OpSys'].str.replace(old_word,new_word)

print(df['OpSys'].value_counts())

```

```

Windows    1125
OS          66
Linux       62
Chrome      27
Mac         21
Android      2
Name: OpSys, dtype: int64

```

In []:

```


plt.figure(figsize=(12, 10))

# First subplot: Countplot
plt.subplot(2, 1, 1)
sns.countplot(x='OpSys', data=df,palette='coolwarm')
plt.title('Count of Laptops by OS')
plt.grid(True, axis='y', color='gray', linestyle='--', linewidth=0.5)

# Second subplot: Bar plot
plt.subplot(2, 1, 2)
sns.barplot(x='OpSys', y='Price_euros', data=df,palette='coolwarm',ci=None)
plt.title('Price of Laptops by different OS')
plt.grid(True, axis='y', color='gray', linestyle='--', linewidth=0.5)

plt.tight_layout()
plt.show()

```



- Its clear that count of windows laptop is greater than all the OS in big margin
- Mac OS price is greater compared to other OS laptops followed by windows with 1200 EU
- Below windows price drops to half of windows OS

Weight

```
In [ ]: df['Weight'] = df['Weight'].str.replace("kg", " ")
df['Weight'] = pd.to_numeric(df['Weight'])
```

- Replacing the "kg" in weight with blank space
- Changing the dtype of weight column to numerical
- To further evaluate

```
In [ ]: plt.figure(figsize=(18, 8))

# First subplot: Countplot
plt.subplot(2, 2, 1)
sns.scatterplot(x='Weight', y='Price_euros', hue='OpSys', data=df, palette='viridis', s=100)
plt.title('Scatterplot of Laptops by Weight')
plt.grid(True, axis='y', color='gray', linestyle='--', linewidth=0.5)

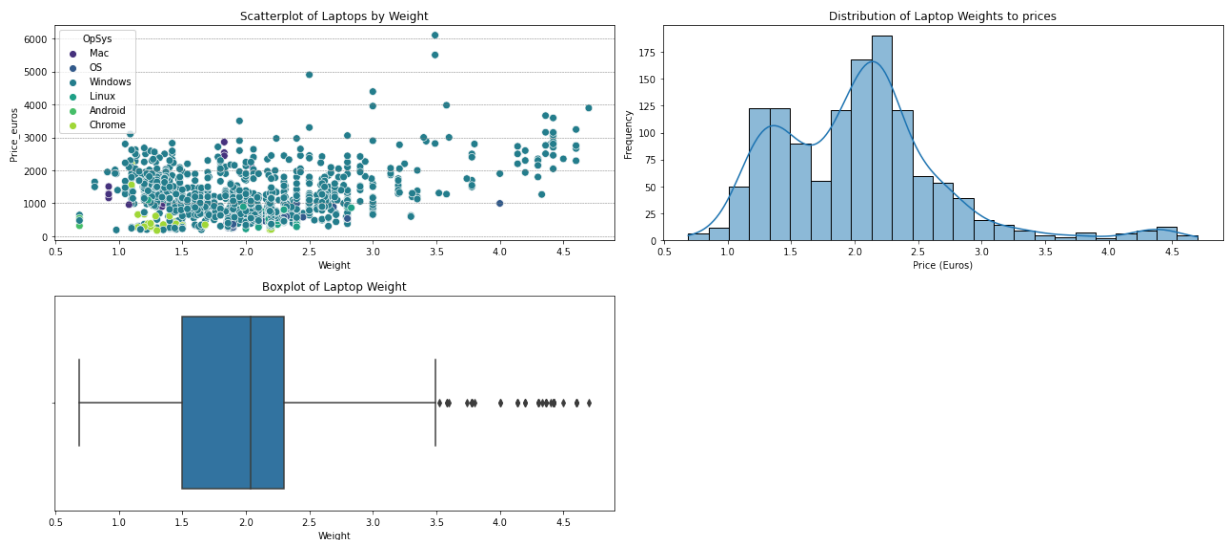
# Second subplot: Histogram
plt.subplot(2, 2, 2)
sns.histplot(x='Weight', data=df, bins=25, kde=True)
plt.title('Distribution of Laptop Weights to prices')
plt.xlabel('Price (Euros)')
```



```
plt.ylabel('Frequency')

# Third subplot : Boxplot
plt.subplot(2, 2, 3)
sns.boxplot(x='Weight', data=df)
plt.title('Boxplot of Laptop Weight')
plt.xlabel('Weight')

plt.tight_layout() # Adjust the layout to prevent overlap
# Show the plot
plt.show()
```



- Windows laptop has occupied in large scale in weight variation
- windows laptops has products 0.7-4.2 kg of weight
- most of the laptops fell between 1.1 to 2.5 kg .
- Median weight of laptop is 2.1kg


Price (euros)

```
In [ ]: # Create a figure with a specified size
plt.figure(figsize=(18, 5))

# First subplot: Histogram
plt.subplot(1, 2, 1)
sns.histplot(x='Price_euros', data=df, bins=30,kde=True)
plt.title('Distribution of Laptop prices')
plt.xlabel('Price (Euros)')
plt.ylabel('Frequency')

# Second subplot : Boxplot
plt.subplot(1, 2, 2)
sns.boxplot(x='Price_euros', data=df)
plt.title('Boxplot of Laptop Weight')
plt.xlabel('Weight')

plt.tight_layout()
plt.show()
```



- 600 to 1500 EU is covers the most of laptop price.
- Above 2900EU is considered as over priced laptops
- Price feature is right skewed

Ram Feature

```
In [ ]: df['Ram'] = df['Ram'].str.replace("GB", " ")
df['Ram'] = pd.to_numeric(df['Ram'])
df['Ram'].value_counts()
```

```
Out[ ]: 8      619
4      375
16     200
6       41
12      25
2       22
32      17
24       3
64       1
Name: Ram, dtype: int64
```


- Replacing the 'GB' to blank
- Changing the dtype to numerical

```
In [ ]: plt.figure(figsize=(18, 5))

# First subplot: Countplot
plt.subplot(1, 2, 1)
sns.countplot(x='Ram', data=df, palette='coolwarm')
plt.title('Count of Laptops by RAM Types')
plt.grid(True, axis='y', color='gray', linestyle='--', linewidth=0.5)

# Second subplot: Bar plot
plt.subplot(1, 2, 2)
sns.barplot(x='Ram', y='Price_euros', data=df, palette='coolwarm', ci=None)
plt.title('Price of Laptops by RAM Types')
plt.grid(True, axis='y', color='gray', linestyle='--', linewidth=0.5) # Add gridlin

plt.tight_layout()
plt.show()
```



- Over 600 laptops has 8GB RAM installed in their laptops
- Second most is the 4GB RAM installed in 380laptops
- Propotional to the RAM Type the price od laptops increases

CPU Feature

```
In [ ]: df['Cpu'].value_counts()
```

```
Out[ ]: Intel Core i5 7200U 2.5GHz      190
Intel Core i7 7700HQ 2.8GHz      146
Intel Core i7 7500U 2.7GHz      134
Intel Core i7 8550U 1.8GHz       73
Intel Core i5 8250U 1.6GHz       72
...
Intel Core M M3-6Y30 0.9GHz       1
AMD A9-Series 9420 2.9GHz         1
Intel Core i3 6006U 2.2GHz         1
AMD A6-Series 7310 2GHz            1
Intel Xeon E3-1535M v6 3.1GHz       1
Name: Cpu, Length: 118, dtype: int64
```

- Every PC builders are PC enthusiast knows that
- whether its high end or low end processor
- Processor generation and frequency is most important than model.
- So we are separating the frequencies mentioned in separate feature

```
In [ ]: df['Cpu_hertz'] = df['Cpu'].str.rsplit(' ', 1).str[1]
df['Cpu_hertz'] = df['Cpu_hertz'].str.replace('GHz', " ")
df['Cpu_hertz'] = pd.to_numeric(df['Cpu_hertz'])
df['Cpu_hertz']
```

```
Out[ ]: 0      2.3
1      1.8
2      2.5
3      2.7
4      3.1
...
1298   2.5
1299   2.5
1300   1.6
1301   2.5
1302   1.6
Name: Cpu_hertz, Length: 1303, dtype: float64
```

- Separating and changing it to Numerical feature

```
In [ ]: df['Cpu_brand'] = df['Cpu'].str.split(' ',1).str[0]
df['Cpu_brand'].value_counts()
```

```
Out[ ]: Intel      1240
AMD          62
Samsung       1
Name: Cpu_brand, dtype: int64
```

```
In [ ]: df = df[~(df['Cpu_brand']=='Samsung')]
```

- Dropping the samsung processor and filter the cpu to two types

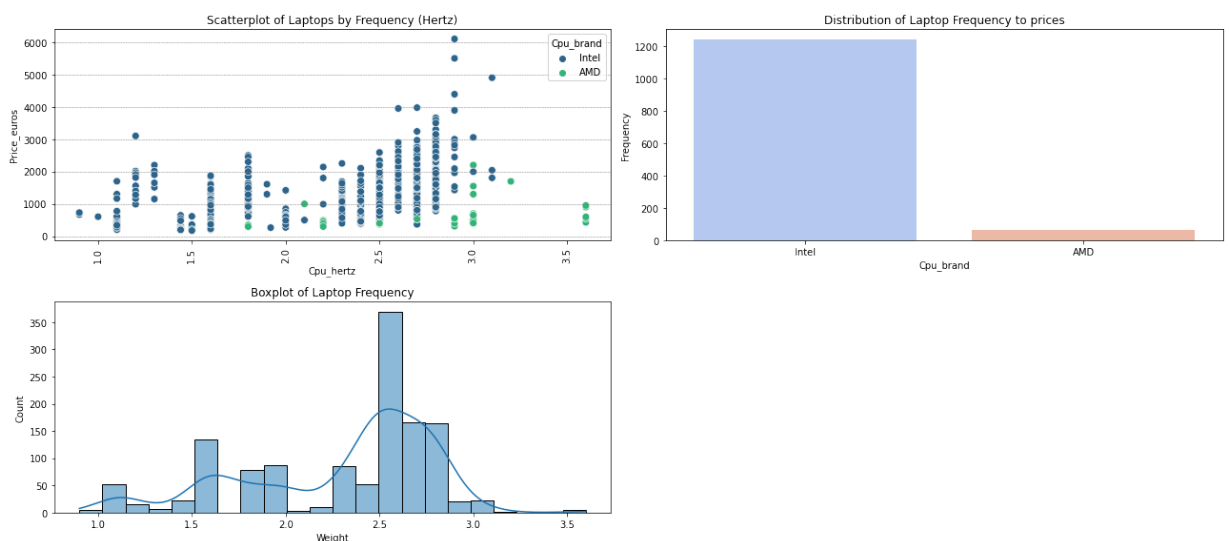
```
In [ ]: plt.figure(figsize=(18, 8))

# First subplot: Countplot
plt.subplot(2, 2, 1)
sns.scatterplot(x='Cpu_hertz', y='Price_euros', hue='Cpu_brand', data=df, palette='vir')
plt.xticks(rotation=90)
plt.title('Scatterplot of Laptops by Frequency (Hertz)')
plt.grid(True, axis='y', color='gray', linestyle='--', linewidth=0.5)

# Second subplot: Histogram
plt.subplot(2, 2, 2)
sns.countplot(x='Cpu_brand', data=df, palette='coolwarm')
plt.title('Distribution of Laptop Frequency to prices')
plt.ylabel('Frequency')

# Third subplot : Boxplot
plt.subplot(2, 2, 3)
sns.histplot(x='Cpu_hertz', data=df, kde=True)
plt.title('Boxplot of Laptop Frequency')
plt.xlabel('Weight')

plt.tight_layout()
plt.show()
```



- Distribution of CPU processor is imbalanced
- More than 90% of samples fell in intel variable
- 2.5 - 2.9 hertz has higher bin of laptops

Resolution - Display

```
In [ ]: df['ScreenResolution'].value_counts()
```

```
Out[ ]: Full HD 1920x1080          507
1366x768          281
IPS Panel Full HD 1920x1080      230
IPS Panel Full HD / Touchscreen 1920x1080    53
Full HD / Touchscreen 1920x1080    47
1600x900          23
Touchscreen 1366x768           16
Quad HD+ / Touchscreen 3200x1800     15
IPS Panel 4K Ultra HD 3840x2160     12
IPS Panel 4K Ultra HD / Touchscreen 3840x2160 11
4K Ultra HD / Touchscreen 3840x2160    10
4K Ultra HD 3840x2160              7
Touchscreen 2560x1440              7
IPS Panel 1366x768                 7
IPS Panel Quad HD+ / Touchscreen 3200x1800   6
IPS Panel Retina Display 2560x1600          6
IPS Panel Retina Display 2304x1440          6
Touchscreen 2256x1504                  6
IPS Panel Touchscreen 2560x1440           5
IPS Panel Retina Display 2880x1800         4
IPS Panel Touchscreen 1920x1200           4
1440x900                                4
IPS Panel 2560x1440                     4
2560x1440                               3
Quad HD+ 3200x1800                      3
1920x1080                               3
Touchscreen 2400x1600                   3
IPS Panel Quad HD+ 2560x1440             3
IPS Panel Touchscreen 1366x768           3
IPS Panel Touchscreen / 4K Ultra HD 3840x2160 2
IPS Panel Full HD 2160x1440              2
IPS Panel Quad HD+ 3200x1800             2
IPS Panel Retina Display 2736x1824       1
IPS Panel Full HD 1920x1200              1
IPS Panel Full HD 2560x1440              1
IPS Panel Full HD 1366x768               1
Touchscreen / Full HD 1920x1080          1
Touchscreen / Quad HD+ 3200x1800         1
Touchscreen / 4K Ultra HD 3840x2160      1
Name: ScreenResolution, dtype: int64
```

- This is to check the informations mentioned in screen resolution
- We can see the touch screen , IPS and resolution dimension would be usefull

```
In [ ]: df['TouchScreen'] = df['ScreenResolution'].apply(lambda element:1 if 'Touchscreen' in element)
df['IPS_panel'] = df['ScreenResolution'].apply(lambda element:1 if 'IPS' in element)
```

- Expanding the features to Touchscreen and IPS panels

```
In [ ]: df['ScreenResolution'] = df['ScreenResolution'].str.extract(r'(\d{3,4}x\d{3,4})')

df[['Screen_width', 'Screen_height']] = df['ScreenResolution'].str.split('x', expand=True)

df.head()
```

Out []:

	Company	TypeName	Inches	ScreenResolution	Cpu	Ram	Memory	Gpu	OpSys	Weight
0	Apple	Ultrabook	13.3	2560x1600	Intel Core i5 2.3GHz	8	128GB SSD	Intel Iris Plus Graphics 640	Mac	1.37
1	Apple	Ultrabook	13.3	1440x900	Intel Core i5 1.8GHz	8	128GB Flash Storage	Intel HD Graphics 6000	Mac	1.34
2	HP	Notebook	15.6	1920x1080	Intel Core i5 7200U 2.5GHz	8	256GB SSD	Intel HD Graphics 620	OS	1.86
3	Apple	Ultrabook	15.4	2880x1800	Intel Core i7 2.7GHz	16	512GB SSD	AMD Radeon Pro 455	Mac	1.83
4	Apple	Ultrabook	13.3	2560x1600	Intel Core i5 3.1GHz	8	256GB SSD	Intel Iris Plus Graphics 650	Mac	1.37

- Extracting the screen resolution only

In []:

```
plt.figure(figsize=(15, 5))


# First subplot: Countplot
plt.subplot(2, 2, 1)
sns.barplot(x='TouchScreen', y='Price_euros', data=df, palette='viridis')
plt.title('Barplot of Laptops by Touch Screen with Price')
plt.grid(True, axis='y', color='gray', linestyle='--', linewidth=0.5)

# Second subplot: Histogram
plt.subplot(2, 2, 2)
sns.countplot(x='TouchScreen', data=df, palette='coolwarm')
plt.title('Countplot of Laptop by Touch Screen')
plt.ylabel('Frequency')

# Third subplot: Barplot
plt.subplot(2, 2, 3)
sns.barplot(x='IPS_panel', y='Price_euros', data=df, palette='viridis')
plt.title('Barplot of Laptops with IPS panel with Price')
plt.grid(True, axis='y', color='gray', linestyle='--', linewidth=0.5)

# Fourth subplot: Countplot
plt.subplot(2, 2, 4)
sns.countplot(x='IPS_panel', data=df, palette='coolwarm')
plt.title('Countplot of Laptop with IPS panel')
plt.ylabel('Frequency')

plt.tight_layout()
plt.show()
```



- Count of Touch screen is low but the cost of touch screen is higher
- Cost price of laptops with IPS panel is higher but count of laptops with IPS panel is 1/3rd.

```
In [ ]: df['PPI'] = (((df['Screen_width']**2+df['Screen_height']**2))**0.5/df['Inches']).ast
df.head()
```

```
Out [ ]: Company  TypeName  Inches  ScreenResolution  Cpu  Ram  Memory  Gpu  OpSys  Weight

0      Apple  Ultrabook   13.3      2560x1600  Intel Core i5 8    128GB SSD  Intel Iris Plus Graphics 640  Mac  1.37

1      Apple  Ultrabook   13.3      1440x900  Intel Core i5 8    128GB Flash Storage  Intel HD Graphics 6000  Mac  1.34

2        HP  Notebook   15.6      1920x1080  Intel Core i5 8    256GB SSD  Intel HD Graphics 620  OS  1.86

3      Apple  Ultrabook   15.4      2880x1800  Intel Core i7 16   512GB SSD  AMD Radeon Pro 455  Mac  1.83

4      Apple  Ultrabook   13.3      2560x1600  Intel Core i5 8    256GB SSD  Intel Iris Plus Graphics 650  Mac  1.37
```

- To cut long story short screen height and width is creating multicollinearity
- To drop either one column is now use
- To utilize the screens resolution we can calculate the PPI with resolution and size of laptop which we are having.

Memory

```
In [ ]: import re

# Function to extract and clean SSD and HDD capacities
def split_storage(storage):
    parts = storage.split(' + ')
    ssd = None
    hdd = None
```

```

for part in parts:
    if 'SSD' in part or 'Flash Storage' in part:
        ssd = part.strip()
    elif 'HDD' in part or 'Hybrid' in part:
        hdd = part.strip()

    return pd.Series([ssd, hdd])

# Apply the function to the Storage column
df[['SSD', 'HDD']] = df['Memory'].apply(split_storage)

# Define a regular expression pattern to remove unwanted text
pattern = r'[A-Za-z\s]+'

# Function to convert TB to GB and add 1000
def convert_to_gb(value):
    if pd.isna(value) and 'TB' in value:
        numeric_value = float(re.sub(pattern, '', value))
        return str(int(numeric_value * 1000))
    return value

# Apply the function to the SSD and HDD columns
df['SSD'] = df['SSD'].apply(convert_to_gb)
df['HDD'] = df['HDD'].apply(convert_to_gb)

# Apply the pattern to the SSD and HDD columns and clean up
df['SSD'] = df['SSD'].str.replace(pattern, '', regex=True).str.strip()
df['HDD'] = df['HDD'].str.replace(pattern, '', regex=True).str.strip()

df['SSD'] = df['SSD'].fillna(0)
df['HDD'] = df['HDD'].fillna(0)

# Print the resulting DataFrame
print(df)

```

	Company		TypeName	Inches	ScreenResolution	\		
0	Apple		Ultrabook	13.3	2560x1600			
1	Apple		Ultrabook	13.3	1440x900			
2	HP		Notebook	15.6	1920x1080			
3	Apple		Ultrabook	15.4	2880x1800			
4	Apple		Ultrabook	13.3	2560x1600			
...			
1298	Lenovo	2 in 1	Convertible	14.0	1920x1080			
1299	Lenovo	2 in 1	Convertible	13.3	3200x1800			
1300	Lenovo		Notebook	14.0	1366x768			
1301	HP		Notebook	15.6	1366x768			
1302	Asus		Notebook	15.6	1366x768			
				Cpu	Ram	Memory \		
0				Intel Core i5 2.3GHz	8	128GB SSD		
1				Intel Core i5 1.8GHz	8	128GB Flash Storage		
2				Intel Core i5 7200U 2.5GHz	8	256GB SSD		
3				Intel Core i7 2.7GHz	16	512GB SSD		
4				Intel Core i5 3.1GHz	8	256GB SSD		
...					
1298				Intel Core i7 6500U 2.5GHz	4	128GB SSD		
1299				Intel Core i7 6500U 2.5GHz	16	512GB SSD		
1300				Intel Celeron Dual Core N3050 1.6GHz	2	64GB Flash Storage		
1301				Intel Core i7 6500U 2.5GHz	6	1TB HDD		
1302				Intel Celeron Dual Core N3050 1.6GHz	4	500GB HDD		
				Gpu	OpSys	Weight	Price_euros	Cpu_hertz \

0	Intel Iris Plus Graphics 640	Mac	1.37	1339.69	2.3
1	Intel HD Graphics 6000	Mac	1.34	898.94	1.8
2	Intel HD Graphics 620	OS	1.86	575.00	2.5
3	AMD Radeon Pro 455	Mac	1.83	2537.45	2.7
4	Intel Iris Plus Graphics 650	Mac	1.37	1803.60	3.1
...
1298	Intel HD Graphics 520	Windows	1.80	638.00	2.5
1299	Intel HD Graphics 520	Windows	1.30	1499.00	2.5
1300	Intel HD Graphics	Windows	1.50	229.00	1.6
1301	AMD Radeon R5 M330	Windows	2.19	764.00	2.5
1302	Intel HD Graphics	Windows	2.20	369.00	1.6

	Cpu_brand	TouchScreen	IPS_panel	Screen_width	Screen_height	\
0	Intel	0	1	2560.0	1600.0	
1	Intel	0	0	1440.0	900.0	
2	Intel	0	0	1920.0	1080.0	
3	Intel	0	1	2880.0	1800.0	
4	Intel	0	1	2560.0	1600.0	
...	
1298	Intel	1	1	1920.0	1080.0	
1299	Intel	1	1	3200.0	1800.0	
1300	Intel	0	0	1366.0	768.0	
1301	Intel	0	0	1366.0	768.0	
1302	Intel	0	0	1366.0	768.0	

	PPI	SSD	HDD
0	226.983005	128	0
1	127.677940	128	0
2	141.211998	256	0
3	220.534624	512	0
4	226.983005	256	0
...
1298	157.350512	128	0
1299	276.053530	512	0
1300	111.935204	64	0
1301	100.454670	0	1000
1302	100.454670	0	500

[1302 rows x 20 columns]

- Separating the SSD and HDD to New features
- imputing 0 were its null

GPU

```
In [ ]: df['Gpu'] = df['Gpu'].str.split(' ', 1).str[0]

df = df[~df['Gpu'].str.contains('ARM')]

df.Gpu.value_counts()
```

```
Out[ ]: Intel      722
Nvidia    400
AMD        180
Name: Gpu, dtype: int64
```


- Graphics distribution

```
In [ ]: # Create a figure with a specified size
plt.figure(figsize=(15, 5))
```

```
# First subplot: Countplot
plt.subplot(1, 2, 1)
sns.barplot(x='Gpu', y='Price_euros', data=df, palette='viridis', ci=None)
plt.xticks(rotation=90)
plt.title('Barplot of Laptops by Gpu')
plt.grid(True, axis='y', color='gray', linestyle='--', linewidth=0.5)

# Second subplot: Histogram
plt.subplot(1, 2, 2)
sns.countplot(x='Gpu', data=df, palette='viridis')
plt.title('Countplot of Laptop by Gpu')
plt.ylabel('Frequency')

plt.tight_layout() # Adjust the layout to prevent overlap
# Show the plot
plt.show()
```



- Count of Intel graphics laptops is high
- Price of Nvidia graphics is greater than Intel and has laptop count greater than AMD

Encoding

In []:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 1302 entries, 0 to 1302
Data columns (total 20 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   Company               1302 non-null   object
 1   TypeName              1302 non-null   object
 2   Inches               1302 non-null   float64
 3   ScreenResolution      1302 non-null   object
 4   Cpu                  1302 non-null   object
 5   Ram                 1302 non-null   int64
 6   Memory              1302 non-null   object
 7   Gpu                 1302 non-null   object
 8   OpSys               1302 non-null   object
 9   Weight             1302 non-null   float64
10   Price_euros         1302 non-null   float64
11   Cpu_hertz          1302 non-null   float64
12   Cpu_brand          1302 non-null   object
13   TouchScreen        1302 non-null   int64
14   IPS_panel          1302 non-null   int64
15   Screen_width       1302 non-null   float64
16   Screen_heigth      1302 non-null   float64
```

```

17  PPI                1302 non-null    float64
18  SSD                1302 non-null    object
19  HDD                1302 non-null    object
dtypes: float64(7), int64(3), object(10)
memory usage: 245.9+ KB

```

```
In [ ]: df.drop(columns=['ScreenResolution', 'Cpu', 'Memory', 'Screen_width', 'Screen_height'], a
```

- Dropping the redundant features
- Because we created clean features from that which we will use in machine learning

```
In [ ]: df['TypeName'] = df['TypeName'].str.replace('2 in 1 ', "")
df['SSD'] = df['SSD'].astype(float)
df['HDD'] = df['HDD'].astype(float)
```

- Changing the dtype to float
- replacing numbers in Typename for encoding

```
In [ ]: # importing Label encoder
from sklearn.preprocessing import LabelEncoder

le = LabelEncoder()
```

```
In [ ]: # Create the 'intel' column
df['intel'] = df.apply(lambda row: 1 if row['Cpu_brand'] == 'Intel' and 'Intel' in row['Cpu_brand'] else 0)

# Create the 'AMD' column
df['AMD'] = df.apply(lambda row: 1 if row['Cpu_brand'] == 'AMD' and 'AMD' in row['Cpu_brand'] else 0)
```

- Creating feature with same processor and graphics installed to check the significance

```
In [ ]: df.head()
```

```
Out[ ]: 
```

	Company	TypeName	Inches	Ram	Gpu	OpSys	Weight	Price_euros	Cpu_hertz	Cpu_brand	Tc
0	Apple	Ultrabook	13.3	8	Intel	Mac	1.37	1339.69	2.3	Intel	
1	Apple	Ultrabook	13.3	8	Intel	Mac	1.34	898.94	1.8	Intel	
2	HP	Notebook	15.6	8	Intel	OS	1.86	575.00	2.5	Intel	
3	Apple	Ultrabook	15.4	16	AMD	Mac	1.83	2537.45	2.7	Intel	
4	Apple	Ultrabook	13.3	8	Intel	Mac	1.37	1803.60	3.1	Intel	

```
In [ ]: for col in df.columns:
        if df[col].dtype == 'object':
            df[col] = le.fit_transform(df[col])

print(df.info())
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 1302 entries, 0 to 1302
```

Data columns (total 17 columns):

#	Column	Non-Null Count	Dtype
0	Company	1302 non-null	int32
1	TypeName	1302 non-null	int32
2	Inches	1302 non-null	float64
3	Ram	1302 non-null	int64
4	Gpu	1302 non-null	int32
5	OpSys	1302 non-null	int32
6	Weight	1302 non-null	float64
7	Price_euros	1302 non-null	float64
8	Cpu_hertz	1302 non-null	float64
9	Cpu_brand	1302 non-null	int32
10	TouchScreen	1302 non-null	int64
11	IPS_panel	1302 non-null	int64
12	PPI	1302 non-null	float64
13	SSD	1302 non-null	float64
14	HDD	1302 non-null	float64
15	intel	1302 non-null	int64
16	AMD	1302 non-null	int64

dtypes: float64(7), int32(5), int64(5)

memory usage: 190.0 KB

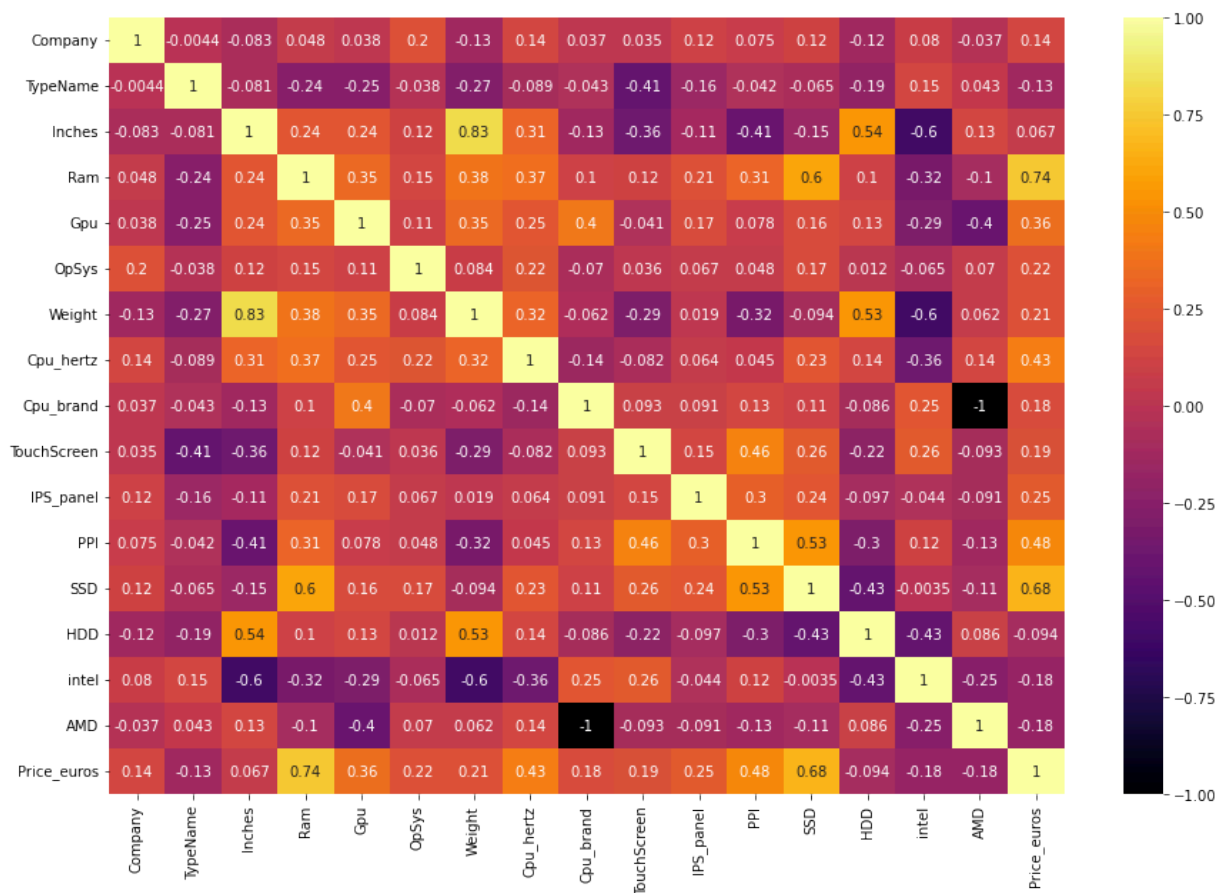
None

- Automating a function to label encode feature for model

```
In [ ]: price_column = df.pop('Price_euros') # Remove 'Salary' column
df.insert(16, 'Price_euros', price_column)
```

- Putting the price feature to one side for easy comparability


```
In [ ]: corr = df.corr()
plt.figure(figsize=(15,10))
sns.heatmap(corr,annot=True,cmap='inferno')
plt.show()
```



```
In [ ]: df.drop(columns=['Inches', 'HDD', 'AMD', 'intel', 'TypeName', 'Weight'], inplace=True)
```

- Checking the Corelation and dropping the features which area reduntant or multicolinear

```
In [ ]: corr = df.corr()
plt.figure(figsize=(16,10))
sns.heatmap(corr,annot=True,cmap='inferno',fmt= ".2f")
plt.show()
```



- Features before modelling

Machine learning model

Data preprocessing

Separating datasets to independent and dependant (Target) variable

```
In [ ]: X = df.drop('Price_euros',axis=1)
        y = np.log(df[['Price_euros']])
```

```
In [ ]: # Importing the libraries for train , test the model
        from sklearn.preprocessing import StandardScaler
        from sklearn.model_selection import train_test_split
        from sklearn.ensemble import RandomForestRegressor
        from sklearn import metrics
```

```
In [ ]: # Splitting data into training and test sets
        X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.20,random_stat
```

```
In [ ]: # Standardizing features
        ss = StandardScaler()
        X_train.loc[:,:] = ss.fit_transform(X_train.loc[:,:])
        X_test.loc[:,:] = ss.transform(X_test.loc[:,:])
```

```
In [ ]: # Training Random Forest Regressor
        rfr = RandomForestRegressor()
        rfr.fit(X_train, y_train)
```

```
Out[ ]: RandomForestRegressor()
```

- Predicting the Target variable for train and test

```
In [ ]: y_train_pred = rfr.predict(X_train)
```

```
In [ ]: y_test_pred = rfr.predict(X_test)
```

- Getting the model score of test data

```
In [ ]: # Evaluating model  
print(f"Model Score: {rfr.score(X_test, y_test)}")
```

Model Score: 0.8556323633303362


```
In [ ]: print('R2 score', metrics.r2_score(y_test, y_test_pred))  
print('MAE', metrics.mean_absolute_error(y_test, y_test_pred))
```

R2 score 0.8556323633303362

MAE 0.1739171746557003

- R2 score and MAE shows that this model is a good fit
- R2 score of 85.4% is good for predicting the laptop price

```
In [ ]: # Plotting predictions  
  
plt.figure(figsize=(12, 8))  
plt.scatter(y_test_pred, y_test)  
plt.xlabel('Predicted Price')  
plt.ylabel('Actual Price')  
plt.title('Price of Actual / Predicted')  
plt.show()
```



- Visual representation of Y actual and Y prediction made from our own model
- The linearity in graph shows that the model has Best fit line

In []:

In []: